

1.0 Cannabis E-Commerce User Behavior & Hybrid Recommendation System

This project analyzes user interactions with products in an e-commerce setting to understand engagement patterns and build a hybrid recommendation system. It combines user-product interaction data with product text information to generate personalized top-N product recommendations. The approach leverages graph embeddings (Node2Vec) for user-product relationships and TF-IDF embeddings for product content, producing a hybrid scoring system that balances user behavior and product similarity.

Dataset:

The dataset consists of two main components. The first contains information about the products available on the website, including attributes such as product name, vendor, product ID, and price. The second captures user behavior and interactions with these products, detailing each user's ID, the products they viewed, the type of interaction (e.g., click, save, add to cart, purchase), the quantity purchased, and any ratings provided.

Goals:

- Explore user-product interactions
- Engineer features capturing engagement
- Build a hybrid recommendation system using text + graph embeddings
- Generate top-5 recommendations for users

2.0 Import Libraries

```
!pip install -q node2vec
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.metrics import classification_report, roc_auc_score
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE
from sklearn.metrics.pairwise import cosine_similarity
import networkx as nx
from node2vec import Node2Vec
from sklearn.preprocessing import normalize
```

3.0 Load Data

```
products_df = pd.read_csv('cannabis-prod-data.csv')
user_behavior_df = pd.read_csv('cannabis-user-data.csv')
```

4.0 Exploratory Data Analysis (EDA)

Data cleaning and exploration.

Here we'll check for missing or inconsistent values in both datasets, validate that product IDs in user behavior match those in the product catalog, explore distributions such as product availability, categories, prices, user event types, rating distributions, purchase quantities, etc.

4a. Quick Overview and Data Inspection

```
print("\n--- Products Info ---")
print(products_df.info())
print(products_df.head())

print("\n--- User Behavior Info ---")
print(user_behavior_df.info())
print(user_behavior_df.head())
```

```
--- Products Info ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 14 columns):
 #   Column          Non-Null Count  Dtype
---  ---
 0   product_id      100 non-null   object
 1   name            100 non-null   object
 2   category        100 non-null   object
 3   strain_type     100 non-null   object
 4   strain          100 non-null   object
 5   vendor          100 non-null   object
 6   THC_pct         100 non-null   float64
 7   CBD_pct         94 non-null    float64
 8   potency_mg      78 non-null    float64
 9   price_usd       100 non-null   float64
10   availability     100 non-null   int64
11   description     89 non-null    object
12   tags            100 non-null   object
13   date_added      100 non-null   object
dtypes: float64(4), int64(1), object(9)
```

```
memory usage: 11.1+ KB
None
product_id      name      category strain_type      strain \
0      p001      OG Kush 3.5g      Flower      Sativa      OG Kush
1      p002      Sour Diesel 1000mg Concentrate Hybrid Sour Diesel
2      p003      Blue Dream One Gram Concentrate Sativa Blue Dream
3      p004      Sour Diesel - 50 pieces Edible Sativa Sour Diesel
4      p005      Sour Diesel - 7g      Flower      Sativa Sour Diesel

      vendor  THC_pct  CBD_pct  potency_mg  price_usd  availability \
0  Herbal Haven    28.5    0.8      NaN    54.36      833
1  CannaWorld     84.0    1.9    750.0    46.02      144
2  Pure Bloom     84.4    0.5    750.0    54.85      954
3  Herbal Haven    17.2    0.3    100.0    23.85      132
4  Happy Trails    22.9    0.5      NaN    94.93      324

      description \
0  OG Kush Sativa flower with energetic, euphoric...
1  Sour Diesel Hybrid concentrate with sleepy, en...
2  Blue Dream Sativa concentrate with uplifting, ...
3  Sour Diesel Sativa edible with cerebral, creat...
4  Sour Diesel Sativa flower with cerebral, eupho...

      tags      date_added
0  ['energetic', 'euphoric', 'uplifting'] 2025-07-29
1  ['sleepy', 'energetic', 'creative', 'pain reli... 2025-06-10
2  ['uplifting', 'focused', 'creative'] 2025-05-17
3  ['cerebral', 'creative', 'uplifting'] 2025-05-04
4  ['cerebral', 'euphoric', 'focused', 'creative'] 2025-03-20

--- User Behavior Info ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29010 entries, 0 to 29009
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---
```

Products Catalog Dataset Overview

The dataset contains 100 cannabis products with details such as product ID, name, category (Flower, Concentrate, Edible, etc.), strain type (Sativa, Indica, Hybrid), strain name, vendor, THC/CBD percentages, potency (in mg where applicable), price (USD), availability count, product description, associated effect/experience tags, and date added. THC content is available for all products, CBD content for most (94%), and potency for a subset (78%). Categories span various consumption formats, and products feature diverse strains like OG Kush, Sour Diesel, and Blue Dream, offered by multiple vendors. Availability values indicate current stock, and textual fields capture marketing descriptions and tags highlighting strain effects such as energetic, creative, and euphoric.

User Behavior Dataset Overview

The dataset contains 29,010 user interactions with products, including user_id, product_id, event type (e.g., click, view, purchase), timestamp, optional product ratings (present for 2,604 events), and purchase quantities. Events are time-stamped at the second level, enabling temporal behavior analysis. The data supports tracking engagement patterns, purchase activity, and user preferences across the product catalog.

```
# check missing values

print("\nMissing values in products:\n", products_df.isna().sum())
print("\nMissing values in user_behavior:\n", user_behavior_df.isna().sum())

Missing values in products:
product_id      0
name            0
category        0
strain_type     0
strain          0
vendor          0
THC_pct         0
CBD_pct         6
potency_mg      22
price_usd       0
availability     0
description     11
tags            0
date_added      0
dtype: int64

Missing values in user_behavior:
user_id         0
product_id      0
event_type      0
timestamp       0
rating          26406
quantity        23163
dtype: int64
```

Missing Values

The products dataframe has 6 rows of missing values in the `CBD_pct` column, 22 rows in the `potency_mg` column, and 11 rows in the `description` column.

The user's behavior dataframe has 26,406 rows without a rating, and 23,163 rows without a quantity observation.

Next, we'll check those columns with any missing values for `CBD_pct`, `potency_mg`, and `description` in the `products_df` dataframe.

```
# check those columns with any missing values CBD_pct, potency_mg, description

products_df[products_df[['CBD_pct', 'potency_mg', 'description']].isna().any(axis = 1)]
```

	product_id	name	category	strain_type	strain	vendor	THC_pct	CBD_pct	potency_mg	price_usd	availability	description	tags	date_added	
	0	p001	OG Kush 3.5g	Flower	Sativa	OG Kush	Herbal Haven	28.5	0.8	NaN	54.36	833	OG Kush Sativa flower with energetic, euphoric...	['energetic', 'euphoric', 'uplifting']	2025-07-29
	4	p005	Sour Diesel - 7g	Flower	Sativa	Sour Diesel	Happy Trails	22.9	0.5	NaN	94.93	324	Sour Diesel Sativa flower with cerebral, eupho...	['cerebral', 'euphoric', 'focused', 'creative']	2025-03-20
	5	p006	CA Gold Rush 1/8 oz	Flower	Indica	CA Gold Rush	Zen Garden	21.0	0.9	NaN	38.01	649	CA Gold Rush Indica flower with body high, rel...	['body high', 'relaxing']	2025-04-22
	9	p010	Wedding Cake - 7g	Flower	Indica	Wedding Cake	Happy Trails	17.1	3.7	NaN	99.45	520	Wedding Cake Indica flower with calming, sleep...	['calming', 'sleepy', 'relaxing', 'body high',...]	2025-06-12
	10	p011	Lemon Cherry 7g	Flower	Sativa	Lemon Cherry	Pure Bloom	31.0	1.1	NaN	92.80	177	Lemon Cherry Sativa flower with creative, focu...	['creative', 'focused', 'energetic']	2025-06-06
	11	p012	OG Kush 7g	Flower	Sativa	OG Kush	Bud Masters	18.3	0.8	NaN	91.17	848	OG Kush Sativa flower with uplifting, focused,...	['uplifting', 'focused', 'creative', 'energeti...]	2025-07-05
	12	p013	Girl Scout Cookies 1000mg	Concentrate	Indica	Girl Scout Cookies	Sunset Wellness	81.0	0.1	1000.0	50.03	170	NaN	['relaxing', 'sedative', 'body high']	2025-07-15
	14	p015	OG Kush 7g	Flower	Hybrid	OG Kush	Herbal Haven	23.1	0.7	NaN	91.38	656	OG Kush Hybrid flower with energetic, upliftin...	['energetic', 'uplifting', 'sleepy']	2025-08-12
	15	p016	Lemon Cherry 3.5g	Flower	Sativa	Lemon Cherry	CannaWorld	20.3	0.3	NaN	57.38	973	Lemon Cherry Sativa flower with cerebral, crea...	['cerebral', 'creative', 'energetic']	2025-04-21
	18	p019	ca gold rush 20 pieces	Edible	Hybrid	CA Gold Rush	Herbal Haven	10.7	NaN	100.0	23.61	0	CA Gold Rush Hybrid edible with energetic, upl...	['energetic', 'uplifting', 'creative', 'sleepy...]	2025-04-03
	19	p020	Blue Dream 50 pieces	Edible	Hybrid	Blue Dream	Emerald Fields	10.0	NaN	50.0	23.39	905	Blue Dream Hybrid edible with relaxing, sleepy...	['relaxing', 'sleepy', 'uplifting', 'calming']	2025-05-04
	20	p021	Pineapple Express - 1/8 oz	Flower	Sativa	Pineapple Express	Zen Garden	15.0	1.4	NaN	41.94	438	Pineapple Express Sativa flower with creative,...	['creative', 'focused', 'uplifting']	2025-07-22
	21	p022	Ice Cream Cake 7g	Flower	Hybrid	Ice Cream Cake	Pure Bloom	21.0	0.8	NaN	108.81	411	Ice Cream Cake Hybrid flower with uplifting, c...	['uplifting', 'creative', 'euphoric']	2025-04-27
	22	p023	Purple Urkle 100g	Edible	Hybrid	Purple Urkle	Happy Trails	16.9	1.3	50.0	25.08	303	NaN	['sleepy', 'calming', 'energetic', 'pain relief']	2025-03-22
	25	p026	Maui Wowie 500mg	Vape Cartridge	Hybrid	Maui Wowie	Elevate Extracts	63.4	0.3	1000.0	46.75	577	NaN	['sleepy', 'creative', 'calming']	2025-05-30
	26	p027	Pineapple Express 500mg	Concentrate	Sativa	Pineapple Express	Sunset Wellness	79.5	NaN	750.0	47.99	597	Pineapple Express Sativa concentrate with focu...	['focused', 'euphoric', 'energetic', 'cerebral...]	2025-02-20
	32	p033	Maui Wowie 1/8 oz	Flower	Hybrid	Maui Wowie	Sunset Wellness	25.0	0.7	NaN	39.90	382	NaN	['uplifting', 'energetic', 'pain relief', 'eup...]	2025-07-05
	36	p037	Girl Scout Cookies 7g	Flower	Indica	Girl Scout Cookies	Bud Masters	16.2	1.6	NaN	94.51	735	Girl Scout Cookies Indica flower with calming,...	['calming', 'body high', 'pain relief', 'sedat...]	2025-02-28
	37	p038	CA Gold Rush - 7g	Flower	Indica	CA Gold Rush	Herbal Haven	16.6	0.0	NaN	99.01	564	CA Gold Rush Indica flower with pain relief, r...	['pain relief', 'relaxing', 'sleepy', 'calming...]	2025-03-03
	46	p047	Jack Herer Eighth	Flower	Indica	Jack Herer	Bud Masters	22.5	1.3	NaN	42.98	400	Jack Herer Indica flower with sleepy, pain rel...	['sleepy', 'pain relief', 'relaxing', 'calming...]	2025-05-15
	48	p049	Maui Wowie 500mg	Concentrate	Sativa	Maui Wowie	Sunset Wellness	77.5	0.9	750.0	46.99	816	NaN	['focused', 'uplifting', 'energetic', 'cerebra...]	2025-04-29
	49	p050	Maui Wowie - 500mg	Concentrate	Indica	Maui Wowie	Happy Trails	82.5	0.5	750.0	47.19	742	NaN	['sedative', 'pain relief', 'sleepy']	2025-07-28
	51	p052	Girl Scout Cookies 7g	Flower	Indica	Girl Scout Cookies	Emerald Fields	18.7	1.1	NaN	99.10	582	Girl Scout Cookies Indica flower with sedative...	['sedative', 'calming', 'pain relief', 'relaxi...]	2025-02-19

Check for duplicate rows

```
# check if there are any duplicate rows

print('There are duplicate rows in product catalog? ', products_df.duplicated().any())

print('\nThere are duplicate rows in user data? ', user_behavior_df.duplicated().any())

There are duplicate rows in product catalog? False

There are duplicate rows in user data? False
```

Merge Datasets

```
# merge datasets on product_id keeping everything on user data

merged_df = pd.merge(user_behavior_df, products_df, on = 'product_id', how = 'left')

print("\n--- Merged Dataset Info ---")
print(merged_df.info())
print(merged_df.head())
```

```
# Column Non-Null Count Dtype
0 user_id 29010 non-null object
1 product_id 29010 non-null object
2 event_type 29010 non-null object
3 timestamp 29010 non-null object
4 rating 2604 non-null float64
5 quantity 5847 non-null float64
6 name 29010 non-null object
7 category 29010 non-null object
8 strain_type 29010 non-null object
9 strain 29010 non-null object
10 vendor 29010 non-null object
11 THC_pct 29010 non-null float64
12 CBD_pct 27385 non-null float64
13 potency_mg 22584 non-null float64
14 price_usd 29010 non-null float64
15 availability 29010 non-null int64
16 description 25740 non-null object
17 tags 29010 non-null object
18 date_added 29010 non-null object
dtypes: float64(6), int64(1), object(12)
memory usage: 4.2+ MB
None
user_id product_id event_type timestamp rating quantity \
0 u0318 p084 view 2025-05-17 00:46:51.853052 NaN NaN
1 u0374 p072 view 2025-05-17 00:50:12.863955 NaN NaN
2 u0711 p087 view 2025-05-17 01:01:42.930708 NaN NaN
3 u0928 p087 view 2025-05-17 01:09:47.040863 NaN NaN
4 u0852 p082 view 2025-05-17 01:10:59.959620 NaN NaN

name category strain_type strain \
0 Pineapple Express 50 pieces Edible Hybrid Pineapple Express
1 Sour Diesel 500mg Concentrate Indica Sour Diesel
2 Girl Scout Cookies 200g Edible Sativa Girl Scout Cookies
3 Girl Scout Cookies 200g Edible Sativa Girl Scout Cookies
4 Jack Herer 1000mg Concentrate Sativa Jack Herer

vendor THC_pct CBD_pct potency_mg price_usd availability \
0 GreenLeaf Co. 16.0 1.1 50.0 26.37 10
1 Herbal Haven 77.6 1.1 750.0 53.76 89
2 Zen Garden 15.1 3.6 50.0 23.32 966
3 Zen Garden 15.1 3.6 50.0 23.32 966
4 Pure Bloom 81.2 3.8 1000.0 47.75 920

description \
0 Pineapple Express Hybrid edible with calming, ...
1 Sour Diesel Indica concentrate with sedative, ...
2 Girl Scout Cookies Sativa edible with cerebral...
3 Girl Scout Cookies Sativa edible with cerebral...
4 Jack Herer Sativa concentrate with uplifting, ...

tags date_added
0 ['calming', 'uplifting', 'energetic', 'euphoric... 2025-06-08
1 ['sedative', 'calming', 'sleepy'] 2025-07-22
2 ['cerebral', 'energetic', 'creative', 'focused'] 2025-03-05
3 ['cerebral', 'energetic', 'creative', 'focused'] 2025-03-05
4 ['uplifting', 'focused'] 2025-08-05
```

4b. Basic Stats:

Unique elements:

```
print(f"\nUnique users: {merged_df['user_id'].nunique()}")
print(f"Unique products: {merged_df['product_id'].nunique()}")
print(f"Unique categories: {merged_df['category'].nunique()}")
print(f"Unique event types: {merged_df['event_type'].nunique()}")

Unique users: 1000
Unique products: 100
Unique categories: 4
Unique event types: 4
```

What categories are there?

```
unique_cat = merged_df['category'].unique()
print("Unique categories: ", unique_cat)

Unique categories:  ['Edible' 'Concentrate' 'Vape Cartridge' 'Flower']
```

What event types are there?

```
unique_event = merged_df['event_type'].unique()
print("Unique event types: ", unique_event)

Unique event types:  ['view' 'click' 'cart' 'buy']
```

Count of each event type:

```
event_counts = merged_df['event_type'].value_counts()
print("\nEvent type counts:\n", event_counts)

Event type counts:
event_type
view      12688
click     9825
cart      3893
buy       2604
Name: count, dtype: int64
```

Min / Max Purchase Quantity

```
min_qty_purchase = merged_df[merged_df['event_type'] == 'buy']['quantity'].min()
max_qty_purchase = merged_df[merged_df['event_type'] == 'buy']['quantity'].max()
```

```
print(f"Minimum purchase quantity: {min_qty_purchase}")
print(f"\nMaximum purchase quantity: {max_qty_purchase}")
```

Minimum purchase quantity: 1.0

Maximum purchase quantity: 10.0

Avg overall purchase quantity and rating

```
avg_qty_purchase = merged_df[merged_df['event_type'] == 'buy']['quantity'].mean()
print(f"Average purchase quantity: {round(avg_qty_purchase, 2)}")
print(f"\nAverage rating: {round(merged_df['rating'].mean(), 2)}")
```

Average purchase quantity: 5.47

Average rating: 3.62

Min / Max Purchase Price and Overall Price

```
min_price_purchase = merged_df[merged_df['event_type'] == 'buy']['price_usd'].min()
max_price_purchase = merged_df[merged_df['event_type'] == 'buy']['price_usd'].max()
min_price = merged_df['price_usd'].min()
max_price = merged_df['price_usd'].max()

print(f"Minimum purchase price: {round(min_price_purchase, 2)}")
print(f"Maximum purchase price: {round(max_price_purchase, 2)}")
print(f"Minimum price: {round(min_price, 2)}")
print(f"Maximum price: {round(max_price, 2)}")
```

Minimum purchase price: 22.84
Maximum purchase price: 109.29
Minimum price: 22.64
Maximum price: 109.29

Min / Max Price Product

```
# get products with the min / max price
min_price_products = merged_df[merged_df['price_usd'] == min_price][['product_id', 'name', 'price_usd', 'category', 'vendor']].drop_duplicates()
max_price_products = merged_df[merged_df['price_usd'] == max_price][['product_id', 'name', 'price_usd', 'category', 'vendor']].drop_duplicates()

min_price_purchase_products = merged_df[merged_df['price_usd'] == min_price_purchase][['product_id', 'name', 'price_usd', 'category', 'vendor']].drop_duplicates()
max_price_purchase_products = merged_df[merged_df['price_usd'] == max_price_purchase][['product_id', 'name', 'price_usd', 'category', 'vendor']].drop_duplicates()
```

```
print("Products Purchased with Minimum Price:")
print(min_price_purchase_products)
```

```
print("\n\nProducts Purchased with Maximum Price:")
print(max_price_purchase_products)
```

```
print("\n\nProducts with Minimum Price:")
print(min_price_products)
```

```
print("\n\nProducts with Maximum Price:")
print(max_price_products)
```

```
Products Purchased with Minimum Price:
product_id  name  price_usd  category  vendor
58      p055  Sour Diesel  100g      22.84  Edible  Sunset Wellness
```

```
Products Purchased with Maximum Price:
product_id  name  price_usd  category  vendor
357      p070  CA Gold Rush  7g      109.29  Flower  Herbal Haven
```

```
Products with Minimum Price:
product_id  name  price_usd  category  vendor
9      p061  Jack Herer  100g      22.64  Edible  Sunset Wellness
```

```
Products with Maximum Price:
product_id  name  price_usd  category  vendor
357      p070  CA Gold Rush  7g      109.29  Flower  Herbal Haven
```

Avg Overall Purchase Price and Avg Overall Price

```
avg_price_purchase = merged_df[merged_df['event_type'] == 'buy']['price_usd'].mean()
avg_price = merged_df['price_usd'].mean()
```

```
print(f"Average purchase price: {round(avg_price_purchase, 2)}")
print(f"\nAverage price: {round(avg_price, 2)}")
```

Average purchase price: 48.38

Average price: 47.12

4c. Distribution Plots

Ratings Distribution

Overall Ratings Distribution

```
rating_colors = sns.color_palette("RdYlGn", 5)

plt.figure(figsize=(8,5))
ax = sns.countplot(x = 'rating', hue = 'rating', data = merged_df, legend = False,
                  order = sorted(merged_df['rating'].dropna().unique()),
                  palette = rating_colors)
plt.title('Rating Distribution')

for p in ax.patches:
    height = p.get_height()
```


```

if height > 0:
    ax.text(
        p.get_x() + p.get_width() / 2, # x-pos: bar start at x-axis then get width divide by two to get center of bar
        height + 1,                    # y-pos: just above bar
        int(height),                    # label as integer
        ha = 'center',                  # horizontal alignment
        va = 'bottom',                  # vertical alignment
    )

for spine in ax.spines.values():
    spine.set_visible(False)

plt.show()

```



Ratings Distribution by Category

```

# group data by category then rating,
data = merged_df.groupby(['category', 'rating']).size().reset_index(name = 'count')

# get unique categories
categories = merged_df['category'].unique()

# create 2x2 subplots
fig, axes = plt.subplots(2, 2, figsize=(12, 8))
axes = axes.flatten() # flatten 2x2 grid to a 1D array for easy looping


# loop through each category and plot
for i, cat in enumerate(categories):
    cat_data = data[data['category'] == cat]
    ax = sns.barplot(x = 'rating', y = 'count', data = cat_data, hue = 'rating',
                    palette = rating_colors, ax = axes[i])
    axes[i].set_title(f'Ratings for {cat}')
    axes[i].set_xlabel('Rating')
    axes[i].set_ylabel('Count')

# add counts on top of bars
for p in ax.patches:
    height = p.get_height()
    if height > 0:
        ax.text(
            p.get_x() + p.get_width() / 2, # x-pos: bar start at x-axis then get width divide by two to get center of bar
            height + 1,                    # y-pos: just above bar
            int(height),                    # label as integer
            ha = 'center',                  # horizontal alignment
            va = 'bottom',                  # vertical alignment
        )

# remove top and right spines
axes[i].spines['top'].set_visible(False)
axes[i].spines['right'].set_visible(False)

plt.tight_layout()
plt.show()

```



Ratings Distribution by Strain Type

```
# group data by strain type then rating
data = merged_df.groupby(['strain_type', 'rating']).size().reset_index(name = 'count')


# get unique strain types
strain_types = merged_df['strain_type'].unique()

for stype in strain_types:
    stype_data = data[data['strain_type'] == stype]
    ax = sns.barplot(x = 'rating', y = 'count', data = stype_data,
                    hue = 'rating', palette = rating_colors)

    ax.set_title(f'Ratings for {stype}')
    ax.set_xlabel('Rating')
    ax.set_ylabel('Count')

    for p in ax.patches:
        height = p.get_height()
        if height > 0:
            ax.text(
                p.get_x() + p.get_width() / 2,
                height + 1,
                int(height),
                ha='center',
                va='bottom' )

    ax.spines['top'].set_visible(False)
    ax.spines['right'].set_visible(False)
plt.tight_layout()
plt.show()
```



Event Type Distribution


Counts of view, click, save, cart, and buy

```
plt.figure(figsize = (8,5))
ax = sns.countplot(x = 'event_type', color = '#41ab5d', data = merged_df)
plt.title('Event Type Distribution')

for p in ax.patches:
    height = p.get_height()
    if height > 0:
        ax.text(
            p.get_x() + p.get_width() / 2,
            height + 1,
            int(height),
            ha = 'center',
            va = 'bottom')

for spine in ax.spines.values():
    spine.set_visible(False)

plt.show()
```


Category Counts for Product

```
plt.figure(figsize=(10,6))
ax = sns.countplot(y = 'category', data = products_df, color = '#4CAF50',
                  order = products_df['category'].value_counts().index)
plt.title('Events by Category')

for p in ax.patches:
    width = p.get_width()
    if width > 0:
        ax.text(
            width + 1,
            p.get_y() + p.get_height() / 2, # middle of bar vertically
            int(width),
            ha = 'left',
            va = 'bottom')

for spine in ax.spines.values():
    spine.set_visible(False)

plt.show()
```




Category Counts for All User Session

```
plt.figure(figsize=(10,6))
ax = sns.countplot(y = 'category', data = merged_df, color = '#74c476',
                  order = merged_df['category'].value_counts().index)
plt.title('Events by Category')

for p in ax.patches:
    width = p.get_width()
    if width > 0:
        ax.text(
            width + 1,
            p.get_y() + p.get_height() / 2, # middle of bar vertically
            int(width),
            ha = 'left',
            va = 'bottom')

for spine in ax.spines.values():
    spine.set_visible(False)

plt.show()
```



```
# get counts per category
category_counts = merged_df['category'].value_counts()


# colors
colors = ['#00441b', '#006d2c', '#238b45', '#41ab5d']

# create pie chart
plt.figure(figsize = (8,8))
wedges, texts, autotexts = plt.pie(
    category_counts,
    labels = category_counts.index,
    autopct = '%1.1f%%',
    startangle = 90,
    colors = colors )

# customize autopct text (percentages inside slices)
for autotext in autotexts:
    autotext.set_color('white')
    autotext.set_fontweight('bold')
    autotext.set_fontsize(12)

# customize slice labels
for label in texts:
    label.set_color('black') # keep labels black
    label.set_fontsize(12)

plt.title('Events by Category')
plt.axis('equal') # make pie circle, not ellipse
plt.show()
```



Price Distribution and Sensitivity


Overall Product Price Distribution

```
# price histogram overall and by category

plt.figure(figsize = (10,6))
ax = sns.histplot(products_df['price_usd'], bins = 30, kde = True, color = "#00441b")
plt.title("Product Price Distribution")
plt.xlabel("Price (USD)")
```

```
for spine in ax.spines.values():
    spine.set_visible(False)

plt.show()
```



Expensive Products

```
# expensive products

expensive_products = products_df[products_df['price_usd'] > 80]

print(expensive_products)
```

	product_id	name	category	strain_type	strain \
4	p005	Sour Diesel - 7g	Flower	Sativa	Sour Diesel
9	p010	Wedding Cake - 7g	Flower	Indica	Wedding Cake
10	p011	Lemon Cherry 7g	Flower	Sativa	Lemon Cherry
11	p012	OG Kush 7g	Flower	Sativa	OG Kush
14	p015	OG Kush 7g	Flower	Hybrid	OG Kush
21	p022	Ice Cream Cake 7g	Flower	Hybrid	Ice Cream Cake
36	p037	Girl Scout Cookies 7g	Flower	Indica	Girl Scout Cookies
37	p038	CA Gold Rush - 7g	Flower	Indica	CA Gold Rush
51	p052	Girl Scout Cookies 7g	Flower	Indica	Girl Scout Cookies
69	p070	CA Gold Rush 7g	Flower	Sativa	CA Gold Rush
78	p079	CA Gold Rush 7g	Flower	Sativa	CA Gold Rush
82	p083	Girl Scout Cookies 7g	Flower	Sativa	Girl Scout Cookies

	vendor	THC_pct	CBD_pct	potency_mg	price_usd	availability \
4	Happy Trails	22.9	0.5	NaN	94.93	324
9	Happy Trails	17.1	3.7	NaN	99.45	520
10	Pure Bloom	31.0	1.1	NaN	92.80	177
11	Bud Masters	18.3	0.8	NaN	91.17	848
14	Herbal Haven	23.1	0.7	NaN	91.38	656
21	Pure Bloom	21.0	0.8	NaN	108.81	411
36	Bud Masters	16.2	1.6	NaN	94.51	735
37	Herbal Haven	16.6	0.0	NaN	99.01	564
51	Emerald Fields	18.7	1.1	NaN	99.10	582
69	Herbal Haven	26.2	1.7	NaN	109.29	417
78	GreenLeaf Co.	22.0	1.7	NaN	103.69	0
82	CannaWorld	16.3	0.8	NaN	101.34	881

	description \
4	Sour Diesel Sativa flower with cerebral, eupho...
9	Wedding Cake Indica flower with calming, sleep...
10	Lemon Cherry Sativa flower with creative, focu...
11	OG Kush Sativa flower with uplifting, focused,...
14	OG Kush Hybrid flower with energetic, upliftin...
21	Ice Cream Cake Hybrid flower with uplifting, c...
36	Girl Scout Cookies Indica flower with calming,...
37	CA Gold Rush Indica flower with pain relief, r...
51	Girl Scout Cookies Indica flower with sedative...
69	CA Gold Rush Sativa flower with creative, upli...
78	NaN
82	Girl Scout Cookies Sativa flower with cerebral...

	tags	date_added
4	['cerebral', 'euphoric', 'focused', 'creative']	2025-03-20
9	['calming', 'sleepy', 'relaxing', 'body high',...	2025-06-12
10	['creative', 'focused', 'energetic']	2025-06-06
11	['uplifting', 'focused', 'creative', 'energeti...	2025-07-05
14	['energetic', 'uplifting', 'sleepy']	2025-08-12
21	['uplifting', 'creative', 'euphoric']	2025-04-27
36	['calming', 'body high', 'pain relief', 'sedat...	2025-02-28
37	['pain relief', 'relaxing', 'sleepy', 'calming...	2025-03-03
51	['sedative', 'calming', 'pain relief', 'relaxi...	2025-02-19
69	['creative', 'uplifting', 'energetic']	2025-07-30
78	['cerebral', 'energetic', 'euphoric', 'uplifti...	2025-04-05
82	['cerebral', 'focused', 'creative']	2025-03-04

```
# show columns: product_id, name, category, price_usd and sort by price
products_sorted = products_df.sort_values(by = 'price_usd', ascending = False)

# Select only these 4 columns
products_sorted[['product_id', 'name', 'category', 'price_usd']]
```


	product_id	name	category	price_usd
69	p070	CA Gold Rush 7g	Flower	109.29
21	p022	Ice Cream Cake 7g	Flower	108.81
78	p079	CA Gold Rush 7g	Flower	103.69
82	p083	Girl Scout Cookies 7g	Flower	101.34
9	p010	Wedding Cake - 7g	Flower	99.45
...	...	...	...	...
59	p060	Lemon Cherry - 100g	Edible	23.02
38	p039	CA Gold Rush 200g	Edible	22.86
54	p055	Sour Diesel 100g	Edible	22.84
98	p099	Green Crack 50 pieces	Edible	22.76
60	p061	Jack Herer 100g	Edible	22.64

100 rows x 4 columns

The top 5 expensive products in our catalog are flowers costing more than \$99.

Price Distribution by Category

```
plt.figure(figsize = (12,8))
sns.boxplot(data = products_df, x = 'category',
            y = 'price_usd', color = '#C8E6C9' )
plt.title("Price Distribution by Category")
plt.xticks(rotation = 45)
plt.show()
```



The box plots show that flowers have the most expensive products with the most variance, whereas concentrates and vape cartridges are about similar in price while edibles are the least expensive products.

Histogram of all four product prices

```
# filter only the four product types
product_types = ['Flower', 'Edible', 'Concentrate', 'Vape Cartridge']
filtered_df = products_df[products_df['category'].isin(product_types)]


# choose a different color for each category
colors = ['#00441b', '#41ab5d', '#238b45', '#006d2c']

# create 2x2 subplots
fig, axes = plt.subplots(2, 2, figsize=(12, 8))
axes = axes.flatten()

# loop through product types and plot histograms
for i, product_type in enumerate(product_types):
    subset = filtered_df[filtered_df['category'] == product_type]
    sns.histplot(
        data = subset,
        x = 'price_usd',
        bins = 20,
        kde = True,
        ax = axes[i],
        color = colors[i],
        edgecolor = 'black'
    )
    axes[i].set_title(f'{product_type} Price Distribution')
    axes[i].set_xlabel('Price (USD)')
```

```
axes[i].set_ylabel('Count')
```

```
# Adjust layout
plt.tight_layout()
plt.show()
```




Purchase Quantity Distribution (Only for Buy Events)

```
purchase_qty = merged_df[merged_df['event_type'] == 'buy']['quantity']
plt.figure(figsize = (10, 8))
ax = sns.histplot(purchase_qty, bins = 20, color = "#74c476")

# remove grid lines
ax.grid(False)

# remove top and right spines
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

plt.title("Purchase Quantity Distribution")
plt.xlabel("Quantity")
plt.show()
```



4d. Aggregations

Conversion Funnel: Click -> Cart/Save/None -> Buy

```
# unique users
total_users = merged_df['user_id'].nunique()

# count all who clicked a product
click_users = merged_df[merged_df['event_type'] == 'click']['user_id'].nunique()

# count all who saved a product
save_users = merged_df[merged_df['event_type'] == 'save']['user_id'].nunique()

# count all who added something to their cart
cart_users_set = set(merged_df[merged_df['event_type'] == 'cart']['user_id'])

# count all who bought something
buy_users_set = set(merged_df[merged_df['event_type'] == 'buy']['user_id'])

# & -> users have to be in both to be a buy user because they have to go through cart
buy_users_set = buy_users_set & cart_users_set

cart_users = len(cart_users_set)
buy_users = len(buy_users_set)
```

Penetration Rates

```
# penetration rates (share of all users)
click_penetration = click_users / total_users * 100
cart_penetration = cart_users / total_users * 100
buy_penetration = buy_users / total_users * 100

print("=== Penetration Rates (of all unique users) ===")
print(f"\nClicked product:    {click_penetration:.2f}%\n")
print(f"Added to cart:         {cart_penetration:.2f}%\n")
print(f"Purchased:             {buy_penetration:.2f}%")
```

```
=== Penetration Rates (of all unique users) ===

Clicked product:    100.00%

Added to cart:      95.90%

Purchased:          90.70%
```

```
# create df for penetration rates
penetration_df = pd.DataFrame({
    "Stage": ["Click", "Cart", "Buy"],
    "Penetration (%)": [click_penetration, cart_penetration, buy_penetration]
})

# set seaborn style
sns.set(style = "whitegrid")

# plot penetration rates
plt.figure(figsize = (10, 5))
ax = sns.barplot(
    x = "Stage", y = "Penetration (%)", hue = "Stage",
    data = penetration_df, palette = "Greens", dodge = False)


# remove grid lines
ax.grid(False)

# remove top and right spines
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

# set title and y-axis limits
plt.title("User Penetration Rates by Funnel Stage")
plt.ylim(0, 100)

# add percentage labels
for i, v in enumerate(penetration_df["Penetration (%)"]):
    plt.text(i, v + 1, f"{v:.1f}%", ha = 'center', fontweight = 'bold')

plt.show()
```




Penetration Rate:

Penetration rate measures the percentage of users who interacted with a product at least once relative to the total user base. In this dataset:

- Clicks: 100% of users clicked on at least one product.
- Added to Cart: 96% of users clicked and then added a product to their cart.
- Buy: 90.7% of users who clicked eventually purchased a product.

These numbers show how widely users engaged with the product at each stage.

Funnel Conversion Rates




Funnel Conversion Rate:

Funnel conversion rate measures the percentage of users moving from one stage of the engagement or purchase funnel to the next. In this dataset:

- Click → Cart: 96% of users who clicked also added the product to their cart.
- Cart → Buy: 94.6% of users who added to cart went on to purchase.

These numbers indicate how efficiently users progress through each stage of the purchase funnel. High conversion rates suggest that most users who start interacting continue through to purchase.

▼ Avg Rating and Total Interactions per Category



```
total_saves
1      0
0      0
2      0
3      0
```

The category with the highest total clicks are edibles with almost 3000 clicks followed by concentrates at 2900, then flowers at 2100, and finally vape cartridges at 1800.

The category with the most buys are concentrates at 800 purchases followed by edibles at 700, then flowers at 650, and finally vape cartridges at 450.

The category with the highest average rating are flowers at 3.68 followed closely by edibles at 3.67. Concentrates had an average rating of 3.62 and vape cartridges had the least average rating with 3.43.

Top 10 Products

```
# get counts per product for 'buy' events
buy_counts = (
    merged_df[merged_df['event_type'] == 'buy']
    .groupby('product_id')
    .size()
    .sort_values(ascending = False) )

# get the 10th highest value
top10_cutoff = buy_counts.iloc[9] # the 10th product's count

# keep all products with count >= cutoff
top_products = buy_counts[buy_counts >= top10_cutoff]

print("\nTop 10 (with ties) purchased products:\n", top_products)
```

```
Top 10 (with ties) purchased products:
product_id
p026      43
p065      41
p067      41
p096      41
p093      40
p051      40
p097      39
p077      38
p050      38
p001      38
p003      38
dtype: int64
```

```
# select col from prod_df to print
products_db = products_df[['product_id', 'name', 'vendor', 'category']].drop_duplicates()

# select the product details from products_db that matches the product id in top products
top_product_details = products_db[products_db['product_id'].isin(top_products.index)]

# add set index to product id and add the purchase counts as a column
top_product_details = top_product_details.set_index('product_id')
top_product_details['purchase_count'] = top_products

# reorder columns
top_product_details = top_product_details[['purchase_count', 'name', 'vendor', 'category']]

# sort by purchase_count descending
top_product_details = top_product_details.sort_values(by = 'purchase_count', ascending = False)

print()
print("\nTop 10 (with ties) purchased products:\n", top_product_details)
```

```
Top 10 (with ties) purchased products:
purchase_count      name      vendor \
product_id
p026      43      Maui Wowie 500mg  Elevate Extracts
p065      41      Jack Herer 1000mg  GreenLeaf Co.
p067      41      Girl Scout Cookies 3.5g  Sunset Wellness
p096      41      Girl Scout Cookies - 500mg  Herbal Haven
p051      40      OG Kush 1000mg  Elevate Extracts
p093      40      Wedding Cake One Gram  Happy Trails
p097      39      Girl Scout Cookies 200g  Elevate Extracts
p001      38      OG Kush 3.5g  Herbal Haven
p003      38      Blue Dream One Gram  Pure Bloom
p050      38      Maui Wowie - 500mg  Happy Trails
p077      38      CA Gold Rush 1/8 oz  Bud Masters

category
product_id
p026      Vape Cartridge
p065      Concentrate
p067      Flower
p096      Concentrate
p051      Concentrate
p093      Vape Cartridge
p097      Edible
p001      Flower
p003      Concentrate
p050      Concentrate
p077      Flower
```

The top purchased product was a vape cartridge by Elevate Extracts of Maui Wowie 500mg with 43 purchases. The highest sold concentrate was the second most purchased product which was between Jack Herer 1000mg by GreenLeaf Co. and Girl Scout Cookies - 500mg by Herbal Haven. The highest sold flower was also tied for the second most purchased product and that was Girl Scout Cookies 3.5g by Sunset Wellness. Concentrates accounted for five of the top 10 (with ties resulting in 11 items), flowers accounted for three, vape cartridges accounted for two, and finally one edible.

Top 3 Products In Each Category


```

top_n = 3 # top N products per category
categories = merged_df['category'].unique()

products_db = merged_df[['product_id', 'name', 'vendor', 'category']].drop_duplicates()

for cat in categories:
    # filter for category with purchase event
    cat_df = merged_df[(merged_df['category'] == cat) & (merged_df['event_type'] == 'buy')]

    # count purchases per product
    buy_counts = cat_df.groupby('product_id').size().sort_values(ascending = False)

    # if category has no purchase print no purchases
    if len(buy_counts) == 0:
        print(f"\nCategory '{cat}' has no purchased products.")
        continue

    # determine cutoff for top N (handle ties)
    cutoff_idx = min(top_n-1, len(buy_counts)-1) # in case fewer than N products
    top_cutoff = buy_counts.iloc[cutoff_idx]

    top_products = buy_counts[buy_counts >= top_cutoff]

    # merge with product info, set index to prod id, add buy count col, reorder col
    top_product_details = products_db[products_db['product_id'].isin(top_products.index)]
    top_product_details = top_product_details.set_index('product_id')
    top_product_details['purchase_count'] = top_products
    top_product_details = top_product_details.reset_index()
    top_product_details = top_product_details[['purchase_count', 'name', 'vendor', 'category']]
    top_product_details = top_product_details.sort_values(by = 'purchase_count', ascending = False)

    print(f"\nTop {top_n} (with ties) purchased products in category '{cat}':\n", top_product_details)

```

```

Top 3 (with ties) purchased products in category 'Edible':
  purchase_count      name      vendor category
0              39  Girl Scout Cookies 200g Elevate Extracts Edible
2              35  Purple Urkle 20 pieces Bud Masters Edible
1              34  Purple Urkle 100g Happy Trails Edible

Top 3 (with ties) purchased products in category 'Concentrate':
  purchase_count      name      vendor category
0              41  Girl Scout Cookies - 500mg Herbal Haven Concentrate
2              41  Jack Herer 1000mg GreenLeaf Co. Concentrate
1              40  OG Kush 1000mg Elevate Extracts Concentrate

Top 3 (with ties) purchased products in category 'Vape Cartridge':
  purchase_count      name      vendor category
1              43  Maui Wowie 500mg Elevate Extracts Vape Cartridge
2              40  Wedding Cake One Gram Happy Trails Vape Cartridge
0              36  Sour Diesel 1000mg CannaWorld Vape Cartridge

Top 3 (with ties) purchased products in category 'Flower':
  purchase_count      name      vendor category
1              41  Girl Scout Cookies 3.5g Sunset Wellness Flower
0              38  CA Gold Rush 1/8 oz Bud Masters Flower
2              38  OG Kush 3.5g Herbal Haven Flower

```

4e. Time Based Analysis

Events per Day

```


# convert timestamp column to datetime
merged_df['timestamp'] = pd.to_datetime(merged_df['timestamp'])

# extract the date
merged_df['date'] = merged_df['timestamp'].dt.date

# group by date
events_per_day = merged_df.groupby('date').size()

plt.figure(figsize = (10,5))
events_per_day.plot(color = "brown")
plt.title("Events per Day")
plt.xlabel("Date")
plt.ylabel("Number of Events")
plt.show()

```



The daily time series shows no trend or seasonality as it appears very stochastic.


Events by Hour

```
# extract hour to own col
merged_df['hour'] = merged_df['timestamp'].dt.hour

plt.figure(figsize = (10,6))
ax = sns.countplot(data = merged_df, x = 'hour', color = "#74c476")

# remove grid lines
ax.grid(False)

plt.title("Events by Hour")
plt.show()
```




Events by Weekday

```
# extract weekday to own col
merged_df['weekday'] = merged_df['timestamp'].dt.day_name()

plt.figure(figsize = (10,6))
ax = sns.countplot(data = merged_df, x = 'weekday',
                  order = ['Sunday', 'Monday', 'Tuesday', 'Wednesday', 'Thursday',
                          'Friday', 'Saturday'],
                  color = "#006d2c")

# remove grid lines
ax.grid(False)

plt.title("Events by Weekday")
plt.show()
```



Avg Time Between First and Last Event per User

```
# sort events by user and time so we can get differences
merged_df = merged_df.sort_values(['user_id', 'timestamp'])

# get the previous event timestamp for each user
merged_df['prev_timestamp'] = merged_df.groupby('user_id')['timestamp'].shift(1)

# get difference (in minutes) between current and previous event
merged_df['time_diff'] = (merged_df['timestamp'] - merged_df['prev_timestamp']).dt.total_seconds() / 60

# flag new session if time gap is more than 30 minutes
merged_df['new_session'] = merged_df['time_diff'] > 30
```

```
# create session ID by cumulatively summing new session flags per user
merged_df['session_id'] = merged_df.groupby('user_id')['new_session'].cumsum()

# calc session duration in minutes (max - min timestamp per session)
session_durations = merged_df.groupby(['user_id', 'session_id']).agg(
    session_length = ('timestamp', lambda x: (x.max() - x.min()).total_seconds() / 60))

# calc avg session length across all users and sessions
avg_session_length = session_durations['session_length'].mean()

print(f"Average session length: {avg_session_length:.2f} minutes")

Average session length: 6.41 minutes
```

4f. User Engagement

Avg Events per User

```
events_per_user = merged_df.groupby('user_id').size()

# avg events per user
avg_events = events_per_user.mean()
print(f"Average events per user: {avg_events:.2f}")

Average events per user: 29.01
```

Distribution of Events per User

```
events_count = events_per_user.value_counts().sort_index()

fig, ax = plt.subplots(figsize = (12,6))


# use the actual event numbers as x
x = events_count.index
y = events_count.values
ax.bar(x, y, color = "#A5D6A7")

ax.axvline(x = avg_events, color = 'green', linestyle='--',
           label=f"Avg = {avg_events:.2f}")
ax.legend()

# remove gridlines
ax.grid(False)

# remove top and right spines
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

# labels and title
ax.set_xlabel("Number of Events per User")
ax.set_ylabel("Number of Users")
ax.set_title("User Activity Distribution")
plt.show()
```



```
buyer_frequency = merged_df[merged_df['event_type']=='buy'].groupby('user_id').size()
print("Average purchase frequency per user:", buyer_frequency.mean())

Average purchase frequency per user: 2.8710033076074972
```

```
# Replace 'u0763' with the user_id you want to inspect
user_id_to_check = 'u0501'

# Filter the DataFrame
user_data = merged_df[merged_df['user_id'] == user_id_to_check]

# Or display nicely in Jupyter Notebook
user_data_filtered = user_data[['product_id', 'event_type', 'timestamp', 'quantity', 'name', 'category']]
user_data_filtered
```

	product_id	event_type	timestamp	quantity	name	category
946	p062	view	2025-05-20 02:34:51.889440	NaN	Blue Dream 500mg	Concentrate
4437	p067	view	2025-05-31 00:26:52.889478	NaN	Girl Scout Cookies 3.5g	Flower
4439	p067	click	2025-05-31 00:29:52.889478	NaN	Girl Scout Cookies 3.5g	Flower
4442	p067	click	2025-05-31 00:33:52.889478	NaN	Girl Scout Cookies 3.5g	Flower
7450	p004	view	2025-06-09 14:58:28.889398	NaN	Sour Diesel - 50 pieces	Edible
11998	p050	view	2025-06-23 18:16:11.889406	NaN	Maui Wowie - 500mg	Concentrate
26092	p029	view	2025-08-06 23:31:52.889448	NaN	Maui Wowie 1000mg	Concentrate
27847	p028	view	2025-08-12 12:17:27.889456	NaN	Ice Cream Cake One Gram	Concentrate
27848	p028	click	2025-08-12 12:19:27.889456	NaN	Ice Cream Cake One Gram	Concentrate
27849	p028	cart	2025-08-12 12:22:27.889456	10.0	Ice Cream Cake One Gram	Concentrate
27854	p028	buy	2025-08-12 12:35:27.889456	10.0	Ice Cream Cake One Gram	Concentrate
28162	p092	view	2025-08-13 10:10:32.889417	NaN	OG Kush 3.5g	Flower
28164	p092	click	2025-08-13 10:14:32.889417	NaN	OG Kush 3.5g	Flower
28166	p092	cart	2025-08-13 10:20:32.889417	9.0	OG Kush 3.5g	Flower
28167	p092	buy	2025-08-13 10:35:32.889417	9.0	OG Kush 3.5g	Flower

```
unique_buyers = merged_df[merged_df['event_type'] == 'buy']['user_id'].nunique()
print(f"Number of unique buyers: {unique_buyers}")
```

Number of unique buyers: 907

```
repeat_buyers = buyer_frequency[buyer_frequency > 1] # keep only users with more than 1 purchase
print("Number of repeat buyers:", len(repeat_buyers))
print("Average purchase frequency among repeat buyers:", repeat_buyers.mean())
```

Number of repeat buyers: 709

Average purchase frequency among repeat buyers: 3.393511988716502

```
repeat_purchase_counts = merged_df[merged_df['event_type']=='buy'].groupby(['user_id', 'product_id']).size()
repeat_purchase_users = repeat_purchase_counts[repeat_purchase_counts > 1].reset_index()['user_id'].nunique()
repeat_purchase_rate = repeat_purchase_users / unique_buyers * 100
print(f"Repeat Purchase Rate per Product: {repeat_purchase_rate:.2f}%")
```

Repeat Purchase Rate per Product: 0.00%

5.0 Feature Engineering

For products:

- extract or encode categorical features
- compute summary stats like avg rating, total sales volume
- generate text embeddings from product names/descriptions for NLP or similarity modeling.

For user behavior:

- aggregate user interactions into features per user-product pair (e.g., total views, last purchase date)
- build user profiles or segmentations (e.g., frequent buyers, browsers)

5a. User-Product Interaction Features

`total_interactions` gives a single engagment score for collaborative filtering or similarity weighting while counting each event separately helps feature richness for ML models

```
# weighted interaction score: click = 1, cart = 2, save = 2, buy = 5
event_weight = {'click':1, 'cart':2, 'save':2, 'buy':5}
merged_df['interaction_score'] = merged_df['event_type'].map(event_weight)
```

```
# aggregate per user-product
user_product_features = merged_df.groupby(['user_id','product_id']).agg(
    total_interactions=('interaction_score','sum'),
    num_clicks=('event_type', lambda x: (x=='click').sum()),
    num_cart=('event_type', lambda x: (x=='cart').sum()),
    num_buy=('event_type', lambda x: (x=='buy').sum()),
    num_save=('event_type', lambda x: (x=='save').sum()),
    avg_rating=('rating','mean'),
    last_interaction=('timestamp','max')
).reset_index()
```

5b. User Features

- Summarize aech user's overall behavior
- Helps segment users (heavy buyers, browsers)
- Avg rating can be used for personalized predictions

```
# calc a col called spent to figure out cumulative user expense during dataset
merged_df['spent'] = merged_df.apply(
    lambda row: row['price_usd'] * row['quantity'] if row['event_type'] == 'buy' and row['quantity'] else 0, axis = 1)
```

```
user_features = merged_df.groupby('user_id').agg(
    total_clicks = ('event_type', lambda x: (x=='click').sum()),
    total_cart = ('event_type', lambda x: (x=='cart').sum()),
    total_buy = ('event_type', lambda x: (x=='buy').sum()),
    total_save = ('event_type', lambda x: (x=='save').sum()),
    avg_rating_given = ('rating','mean'),
    total_spent = ('spent','sum') ).reset_index()
```

5c. Product Features

- Summarizing product characteristics for modeling and content based recommendations
- Total buys and saves can help indicate popularity

```
# product-level aggregated features
product_features = merged_df.groupby('product_id').agg(
    category = ('category', 'first'),
    strain_type = ('strain_type', 'first'),
    avg_price = ('price_usd', 'mean'),
    avg_thc = ('THC_pct', 'mean'),
    avg_cbd = ('CBD_pct', 'mean'),
    avg_rating = ('rating', 'mean'),
    total_buys = ('event_type', lambda x: (x=='buy').sum()),
    total_saves = ('event_type', lambda x: (x=='save').sum()) ).reset_index()
```

5d. Time-Based Features

- Captures temporal patterns
- Recency can improve recommendation ranking (RFM-style features)

```
# interaction recency: days since last interaction per user-product
current_date = datetime.now()
user_product_features['days_since_last'] = (
    current_date - user_product_features['last_interaction']).dt.days
```

5e. Text Features (for NLP / Embeddings)

- For building text-based similarity
- `product_text_matrix` can be used to compute cosine similarity between products

```
# combine product name, strain_type, category, description, and tags for embeddings
products_df['text_feature'] = products_df[
    ['name', 'strain_type', 'category', 'description', 'tags']
].fillna('').agg(' '.join, axis = 1)

# semantic similarity
# turn product text into numeric vectors; rows = prod, col = top 500 terms
# values are importance of term to prod
tfidf = TfidfVectorizer(max_features = 500)
product_text_matrix = tfidf.fit_transform(products_df['text_feature'])

# store feature names for interpretation if needed
product_feature_names = tfidf.get_feature_names_out()
```

5f. Collaborative Filtering

Algorithm predicts the missing pieces based on what similar users or similar products are doing

```
user_item_matrix = user_product_features.pivot(
    index = 'user_id',
    columns = 'product_id',
    values = 'total_interactions').fillna(0)
```

6.0 Modeling

Purchase Propensity:

- predict likelihood of purchase given prior behavior and product attributes

We aim to model users' browsing behaviors to forecast which users are likely to convert into buyers. The ultimate goal is to identify high-propensity buyers and target them with tailored recommendations that drive purchases.

When evaluating our models, we will focus on the recall metric which measures how well the model identifies all the positive cases. Recall answers the question: out of all users who would actually make a purchase, what fraction does our model correctly identify? Missing potential buyers carries a much higher opportunity cost than incorrectly targeting users with personalized recommendations. Additionally, false positives are less harmful in this context, since the recommendations are intended to generate leads and encourage conversions.

6a. Purchase Propensity: Classification Model with Logistic Regression, Random Forest, and XGBoost

Prepare the models

Check for class imbalance and if there are we need to correct them otherwise there will be bias.

```
# count event distribution
event_counts = merged_df['event_type'].value_counts(normalize = True)

print(event_counts)
```

```
event_type
view    0.437366
click   0.338676
cart    0.134195
buy     0.089762
Name: proportion, dtype: float64
```

Since `buy` is less than 10% (0.0898), it indicates there is a class imbalance.

Create target variable

For classification, we'll use (1 = buy) and (0 = not buy)

```
# binary target column, True = 1, False = 0
merged_df['is_buy'] = (merged_df['event_type'] == 'buy').astype(int)
```

Select features

- numerical feats: price, thc_pct, cbd_pct
- categorical feats encoded: category, strain_type
- behavioral feats: user-level aggregates

```
# simple features first
X = merged_df[['price_usd', 'THC_pct', 'category', 'strain_type']]

# target
y = merged_df['is_buy']

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size = 0.2, random_state = 42, stratify = y)
```

Scale features and handle categorical encoding using (ColumnTransformer)

```
num_features = ['price_usd', 'THC_pct']
cat_features = ['category', 'strain_type']

preprocessor = ColumnTransformer([
    ('num', StandardScaler(), num_features),
    ('cat', OneHotEncoder(handle_unknown = 'ignore'), cat_features) ])
```

Logistic Regression

```
log_reg_model = Pipeline([
    ('preprocessor', preprocessor),
    ('clf', LogisticRegression(max_iter = 1000, class_weight = 'balanced')) ])
```

```
log_reg_model.fit(X_train, y_train)
print("Pipeline fitted successfully")
```

Pipeline fitted successfully

Random Forest

```
rf_model = Pipeline([
    ('preprocessor', preprocessor),
    ('clf', RandomForestClassifier(n_estimators = 200, class_weight = 'balanced'))
])
```

```
rf_model.fit(X_train, y_train)
print("Pipeline fitted successfully")
```

Pipeline fitted successfully

XGBoost Classifier

```
xgb_model = Pipeline([
    ('preprocessor', preprocessor),
    ('clf', XGBClassifier(
        n_estimators = 300,
        learning_rate = 0.1,
        max_depth = 6,
        scale_pos_weight =
            (y_train.value_counts()[0] / y_train.value_counts()[1]) # handle imbalance
    ))
])
```

```
xgb_model.fit(X_train, y_train)
print("Pipeline fitted successfully")
```

Pipeline fitted successfully

6b. Preliminary Classification Evaluation

```
for model, name in [(log_reg_model, "Logistic Regression"),
                    (rf_model, "Random Forest"),
                    (xgb_model, "XGBoost")]:
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[:,1]

    print(f"--- {name} ---")
    print(classification_report(y_test, y_pred))
    print("ROC-AUC:", roc_auc_score(y_test, y_prob))
    print()
```

```
--- Logistic Regression ---
      precision    recall  f1-score   support

     0       0.92      0.50      0.65      5281
     1       0.10      0.55      0.17       521

 accuracy      0.51
 macro avg      0.51
weighted avg      0.85
```

ROC-AUC: 0.5283257511355124

```
--- Random Forest ---
```

	precision	recall	f1-score	support
0	0.92	0.40	0.56	5281
1	0.10	0.65	0.17	521
accuracy			0.42	5802
macro avg	0.51	0.53	0.36	5802
weighted avg	0.85	0.42	0.52	5802

ROC-AUC: 0.5331825495447592

	precision	recall	f1-score	support
0	0.92	0.37	0.53	5281
1	0.10	0.68	0.17	521
accuracy			0.40	5802
macro avg	0.51	0.52	0.35	5802
weighted avg	0.85	0.40	0.50	5802

ROC-AUC: 0.5316854940446702

1. What the results tell us

- Class 0 (non-purchase) dominates. Precision is very high (0.92), meaning when the model predicts "no purchase," it's almost always right.
- Class 1 (purchase) suffers. Precision is 0.10 (lots of false positives), but recall is surprisingly high (0.55–0.68). That means the models are catching over half of actual purchases, but they're also mislabeling many non-purchases as purchases.
- Accuracy (~40–51%) isn't meaningful here because the dataset is imbalanced.
- ROC-AUC (~0.53) is barely above random guessing (0.50), so the models don't separate the classes well.

2. Why this happens

- Huge class imbalance (5,281 vs. 521). Purchases are only ~9% of the data.
- Likely features aren't strong enough yet to distinguish purchase from non-purchase.
- Logistic Regression struggles with nonlinearities, while RF and XGBoost found more recall but at the cost of precision.

3. Next steps to improve

- Handle imbalance better
 - Instead of just class_weight, try SMOTE (Synthetic Minority Oversampling Technique) to upsample purchase events.
 - Alternatively, undersample the majority class (non-purchases).
- Feature engineering
 - Add more user-level history: e.g., of prior clicks before purchase
 - Add product attributes: THC%, category, strain type, rating, etc.
 - Add interaction features: user-product frequency, average session activity.
- Model tuning
 - Logistic regression may not be powerful enough here → Random Forest & XGBoost need hyperparameter tuning (grid search or randomized search).
 - Use probability thresholds: Right now default is 0.5. If you lower it (e.g., 0.3), you'll catch more purchases (boost recall).
- Evaluation metric
 - Accuracy isn't useful here. Use:
 - Precision-Recall AUC (better for imbalance)
 - F2-score (weights recall higher, which is useful if missing purchases is costly).

6c. SMOTE & Hyperparameter Tuning

- Balance the classes with SMOTE so the models get enough exposure to the minority class.
- Run Logistic Regression, Random Forest, and XGBoost again, but this time we'll include hyperparameter tuning (via GridSearchCV).
- Compare results (precision, recall, F1, ROC-AUC) to see whether balancing and tuning improve the models.

> Applying SMOTE

↳ 8 cells hidden

✓ Hyperparameter tuning: re-run models with gridsearch

✓ Logistic Regression

```
param_grid_lr = {'C': [0.01, 0.1, 1, 10], 'solver': ['liblinear']}
grid_lr = GridSearchCV(LogisticRegression(class_weight = 'balanced', max_iter = 1000),
                       param_grid_lr, scoring = 'roc_auc', cv = 5)
grid_lr.fit(X_train_res, y_train_res);
```

✓ Random Forest

```
param_grid_rf = {'n_estimators': [100, 200],
                 'max_depth': [None, 10, 20]}
grid_rf = GridSearchCV(RandomForestClassifier(class_weight = 'balanced'),
                       param_grid_rf, scoring = 'roc_auc', cv = 5)
grid_rf.fit(X_train_res, y_train_res);
```

XGBoost

```
param_grid_xgb = {'n_estimators': [100, 200],
                  'max_depth': [3, 6, 10],
                  'learning_rate': [0.01, 0.1]}

grid_xgb = GridSearchCV(XGBClassifier(eval_metric = 'logloss'),
                        param_grid_xgb, scoring = 'roc_auc', cv = 5)

grid_xgb.fit(X_train_res, y_train_res);
```

Evaluation

```
for name, model in zip(["Logistic Regression", "Random Forest", "XGBoost"],
                       [grid_lr, grid_rf, grid_xgb]):
    y_pred = model.predict(X_test_enc)
    y_prob = model.predict_proba(X_test_enc)[:,1]

    print(f"\n--- {name} ---")
    print(classification_report(y_test, y_pred))
    print("ROC-AUC:", roc_auc_score(y_test, y_prob))
```

```
--- Logistic Regression ---
              precision    recall  f1-score   support

     0       0.92      0.58      0.71      6602
     1       0.10      0.49      0.17       651

 accuracy          0.57      7253
 macro avg          0.51      0.54      0.44      7253
 weighted avg       0.85      0.57      0.66      7253

ROC-AUC: 0.5318372778160135

--- Random Forest ---
              precision    recall  f1-score   support

     0       0.93      0.40      0.56      6602
     1       0.10      0.67      0.17       651

 accuracy          0.43      7253
 macro avg          0.51      0.54      0.37      7253
 weighted avg       0.85      0.43      0.53      7253

ROC-AUC: 0.5370628506652781

--- XGBoost ---
              precision    recall  f1-score   support

     0       0.93      0.40      0.56      6602
     1       0.10      0.67      0.17       651

 accuracy          0.43      7253
 macro avg          0.51      0.54      0.37      7253
 weighted avg       0.85      0.43      0.53      7253

ROC-AUC: 0.5369374406396423
```

6d. Summary

After applying SMOTE, accuracy increased slightly for Logistic Regression and XGBoost. However, this improvement is not particularly meaningful because balancing the dataset makes accuracy less informative—random guessing would already achieve around 50% accuracy on a balanced dataset.

The key gain from SMOTE is improved recall for the minority class, which is its primary purpose. SMOTE increased the model's sensitivity to minority samples, but the accuracy gains mostly reflect the artificial balance rather than true generalization.

The ROC-AUC (~0.53) remained largely unchanged, suggesting that the models are not learning strong decision boundaries. Oversampling interpolates existing samples but doesn't add new information, indicating a need for stronger features, more domain-driven signals, or alternative resampling/modeling strategies.

7.0 Recommendation System

7a. Graph System

Create user-product undirected bipartite graph for interaction-based embeddings `user_embeddings` and `product_embeddings`

```
# initiate the graph
G = nx.Graph()

# add user nodes
G.add_nodes_from(user_features['user_id'], bipartite = 'user')

# add product nodes
G.add_nodes_from(products_df['product_id'], bipartite = 'product')

# add edges with weights from user_product_features
for _, row in user_product_features.iterrows():
    user = row['user_id']
    product = row['product_id']
    weight = row['total_interactions']
    if weight > 0:
        G.add_edge(user, product, weight = weight)

# create node embeddings using node2vec
# node2vec learns embeddings for both users and products capturing graph structure
node2vec = Node2Vec(G, dimensions = 64, walk_length = 10, num_walks = 50, workers = 2)

model = node2vec.fit(window = 5, min_count = 1, batch_words = 4)

# extract embeddings for products and users as dictionary
user_embeddings = {node: model.wv[node] for node in user_features['user_id'] if node in model.wv}
```



```
product_embeddings = {node: model.wv[node] for node in products_df['product_id'] if node in model.wv}
```

Next, we'll combine product text embeddings (`product_text_matrix`) with graph embeddings (`product_embeddings`) to create a hybrid product embedding to use for cosine similarity between products to create heatmap later to visualize which products are similar.

```
# convert sparse tfidf matrix from engineered feature to dense array
# so only shows product
product_text_dense = product_text_matrix.toarray()

# get product id list to loop with prod embedding dictionary
product_ids = products_df['product_id'].tolist()

# get graph embeddings for each prod in prod embedding dictionary
# if prod didn't get a vector use zero vector
graph_emb_matrix = np.array([product_embeddings.get(p, np.zeros(64)) for p in product_ids])

# concat graph + text embeddings for hybrid representation
# horizontal stack add col to each row so all feats in one row
hybrid_product_embeddings = np.hstack([product_text_dense, graph_emb_matrix])
```

Cosine similarity between products using hybrid product embeddings

```
# gives square matrix where i,j is how similar prod i is to j
product_similarity = cosine_similarity(hybrid_product_embeddings)
```

7b. Hybrid Recommendation Approach

We're going to be using a recommendation system that includes both text based scoring and graph based scoring. The text based scoring will use cosine similarity while the graph based scoring will be using a dot product of normalized embeddings (user @ product).

We're also going to be employing a weighted average for the text based scoring based on the user's interaction with that product. Weights will use the total interaction score we engineered in section 5a under User-Product Interaction. The weights are as follows:

- click = 1
- cart = 2
- save = 2
- buy = 5

These summed weights for each product the user interacted with will be applied to each product embedding culminating in a weighted average for that product which will be associated with that user profile. The user's embeddings will contain only the products the user interacted, then the weights are applied where purchases have most influence. The resulting user's embedding with the weight's applied will be their preferences in the same vector space as the product embeddings, 500 dimensions.

Then we apply cosine similarity between the user's preference profile and each product embedding creating a score for each product ranging from -1 to 1. -1 meaning the vectors point in the opposite direction indicating maximum dissimilarity, 0 being orthogonal vectors meaning no relation, and 1 meaning the vectors point in the same direction indicating maximum similarity.

If the user did not interact with any products we recommend top popular products that users have interacted with.

Note, we used weighted average because it captures more nuance and reflects personalized strength and preferences than a simple average which would imply equal weights to interactions. Sum and Max weren't used because they can apply bias with their magnitude. Sum would reflect the total interaction volume and max could highlight key product features but neglect frequency of interaction.

Hyperparameters for top and weights for the final score. Alpha will be applied to the graph score and (1-alpha) will be applied to the text score. We also need to instantiate the recommendations dictionary to record users recommendations.

```
top_n = 5
alpha = 0.6
recommendations = {}
```

Normalize product embeddings for stable scoring so it ranges from 0 to 1

```
graph_emb_matrix = normalize(graph_emb_matrix)
product_text_dense = normalize(product_text_dense)
```

We're going to loop each user and for every user we'll compute the graph based scores and then text based scores for every product interaction then combine those two scores weighted with alpha and (1-alpha) then record the top 5 in the recommendations dictionary.

```
for user in user_features['user_id']:
    if user in user_embeddings:
        # first compute graph based scores
        # for this user take the graph embedding vector make it a row and normalize
        user_emb_graph = normalize(user_embeddings[user]).reshape(1, -1)

        # graph based scores and flatten to 1d array so each element is a score for one prod
        scores_graph = graph_emb_matrix @ user_emb_graph.T
        scores_graph = scores_graph.ravel()

        # then compute text based scores
        # get all prods user has interacted with to create their user embedding
        user_products = user_product_features[user_product_features['user_id'] == user]['product_id']
        # check if user has interacted with any at all if no intearction, text score = 0
        if not user_products.empty:
            # convert prod ids to indices to access their embeddings
            product_idxs = [product_ids.index(pid) for pid in user_products if pid in product_ids]
            # if user has interacted prods, take text based embeddings of those prods
            if product_idxs:
                user_prod_text = product_text_dense[product_idxs]
                # get total interactions as weight for each prod the user interacted with so it matches their embedding
                weights = user_product_features[
                    (user_product_features['user_id'] == user) &
                    (user_product_features['product_id'].isin(product_ids))
                ]['total_interactions'].values

            # weighted average
            weighted_profile = np.average(user_prod_text, axis = 0, weights = weights)
```

```

user_text_profile = weighted_profile.reshape(1, -1)

# find cosine similarity between user profile and product embeddings
scores_text = cosine_similarity(product_text_dense, user_text_profile).ravel()

else:
    scores_text = np.zeros(len(product_ids))

# combine weighted scores
scores = alpha * scores_graph + (1 - alpha) * scores_text

# exclude interacted prods, sort score by descending and get top_n
interacted = set(user_products)
ranked = np.argsort(scores)[::-1]
recommended_products = [product_ids[i] for i in ranked if product_ids[i] not in interacted] [:top_n]

else:
    # recommend top popular products if not interacted with any product
    popular_products = merged_df.groupby('product_id')['interaction_score'].sum().sort_values(ascending = False).index.tolist()
    recommended_products = [p for p in popular_products if p not in interacted] [:top_n]

recommendations[user] = recommended_products

```

```

# show top-N for first 5 users
for u in list(recommendations.keys())[:5]:
    print(f"User {u}: Recommended Products {recommendations[u]}")

```

```

User u0001: Recommended Products ['p017', 'p061', 'p055', 'p069', 'p037']
User u0002: Recommended Products ['p043', 'p040', 'p078', 'p044', 'p027']
User u0003: Recommended Products ['p079', 'p019', 'p044', 'p004', 'p077']
User u0004: Recommended Products ['p085', 'p079', 'p003', 'p046', 'p019']
User u0005: Recommended Products ['p094', 'p048', 'p078', 'p072', 'p004']

```

```

# create mapping from product_id to product name
product_id_to_name = dict(zip(products_df['product_id'], products_df['name']))

# print recommendations as product_name: product_id
for user in list(recommendations.keys())[:5]:
    recommended_dicts = [{product_id_to_name[pid]: pid} for pid in recommendations[user]]
    print(f"User {user}: Recommended Products {recommended_dicts}")

```

```

User u0001: Recommended Products [{'OG Kush 100g': 'p017'}, {'Jack Herer 100g': 'p061'}, {'Sour Diesel 100g': 'p055'}, {'Purple Urkle 3.5g': 'p069'}, {'Girl Scout Cookies 7g': 'p037'}]
User u0002: Recommended Products [{'Maui Wowie 500mg': 'p043'}, {'Maui Wowie One Gram': 'p040'}, {'Sour Diesel - 100g': 'p078'}, {'Sour Diesel 500mg': 'p044'}, {'Pineapple Express 100g': 'p027'}]
User u0003: Recommended Products [{'CA Gold Rush 7g': 'p079'}, {'CA Gold Rush 20 pieces': 'p019'}, {'Sour Diesel 500mg': 'p044'}, {'Sour Diesel - 50 pieces': 'p004'}, {'CA Gold Rush 100mg': 'p077'}]
User u0004: Recommended Products [{'Blue Dream - 1000mg': 'p085'}, {'CA Gold Rush 7g': 'p079'}, {'Blue Dream One Gram': 'p003'}, {'Blue Dream 1000mg': 'p046'}, {'CA Gold Rush 100mg': 'p019'}]
User u0005: Recommended Products [{'Sour Diesel 1000mg': 'p094'}, {'Sour Diesel 1000mg': 'p048'}, {'Sour Diesel - 100g': 'p078'}, {'Sour Diesel 500mg': 'p072'}, {'Sour Diesel 100mg': 'p004'}]

```

Create a fct where we use userid and recommendation dic to call recommended prods in userid list

```

def user_recs(user_input, recs_dict):
    """
    Fetch recommended products for a specific user or list of users.

    Args:
        user_input: single user_id or list of user_ids
        recs_dict: dictionary from recommendations() function

    Returns:
        If single user_id: list of recommended products
        If list of user_ids: dict {user_id: recommended_products}
    """

    if isinstance(user_input, list):
        return {user: recs_dict.get(user, []) for user in user_input}
    else:
        return recs_dict.get(user_input, [])

```

```
user_recs('u0001', recommendations)
```

```
['p017', 'p061', 'p055', 'p069', 'p037']
```

7c. Distribution of Interaction Scores

```

# aggregate total interaction per product
product_interx = merged_df.groupby('product_id')['interaction_score'].sum().reset_index()
product_interx = product_interx.merge(products_df[['product_id', 'name']], on = 'product_id')
plotdata = product_interx.sort_values('interaction_score', ascending = False)


# plot

plt.figure(figsize = (12, 6))
ax = sns.barplot(x = 'name', y = 'interaction_score', color = '#4CAF50',
                data = plotdata.head(10))

# remove grid lines
ax.grid(False)

plt.xticks(rotation = 45, ha = 'right')
plt.title('Top 10 Products by Total Interaction Score')
plt.ylabel('Total Interaction Score')
plt.xlabel('Product Name')
plt.show()

```



```
product_intrx.sort_values('interaction_score', ascending = False).head(10)
```

	product_id	interaction_score	name
64	p085	428.0	Jack Herer 1000mg
25	p026	426.0	Maui Wowie 500mg
95	p096	423.0	Girl Scout Cookies - 500mg
92	p093	404.0	Wedding Cake One Gram
66	p067	404.0	Girl Scout Cookies 3.5g
96	p097	401.0	Girl Scout Cookies 200g
50	p051	398.0	OG Kush 1000mg
49	p050	396.0	Maui Wowie - 500mg
47	p048	388.0	Sour Diesel 1000mg
84	p085	382.0	Blue Dream - 1000mg

```
product_similarity_df = pd.DataFrame(product_similarity, index = product_ids, columns = product_ids)
product_similarity_df.head()
```

	p001	p002	p003	p004	p005	p006	p007	p008	p009	p010	...	p091	p092	p093	p094	p095	p096	p097	p098	p099
p001	1.000000	0.047067	0.208653	0.220992	0.190093	0.134202	0.134735	0.341860	0.166958	0.195686	...	0.236359	0.274133	0.189225	0.168073	0.151719	0.285036	0.228863	0.124401	0.328918
p002	0.047067	1.000000	0.268223	0.284375	0.261650	0.192289	0.205424	0.088946	-0.034339	0.227722	...	0.086643	0.061663	0.271928	0.234152	0.180021	0.259025	0.187495	0.161071	0.078857
p003	0.208653	0.268223	1.000000	0.239472	0.119553	0.222225	0.147968	0.156499	0.233040	0.268964	...	0.237340	0.268239	0.131886	0.246808	0.082305	0.230183	0.220844	0.268238	0.247186
p004	0.220992	0.284375	0.239472	1.000000	0.275481	0.230977	0.229604	0.138344	0.108995	0.150275	...	0.234509	0.234118	0.271903	0.143709	0.172044	0.312642	0.200457	0.197711	0.190636
p005	0.190093	0.261650	0.119553	0.275481	1.000000	0.150909	0.211261	0.215222	0.167669	0.042093	...	0.162815	0.265833	0.202440	0.424799	0.118384	0.307772	0.153387	0.207558	0.108316


5 rows x 100 columns

7d. Heatmap of Top 10 Interacted Product Similarities

```
# select top 10 interacted products
top_10products_df = product_intrx.sort_values('interaction_score', ascending = False).head(10)
top_10product_ids = top_10products_df['product_id'].tolist()
top_10product_names = top_10products_df['name'].tolist()

# subset the similarity matrix for these top 10 interacted products
top_similarity_matrix = product_similarity_df.loc[top_10product_ids, top_10product_ids]

# plot heatmap
plt.figure(figsize = (12, 10))
sns.heatmap(top_similarity_matrix,
            xticklabels = top_10product_ids,
            yticklabels = top_10product_ids,
            cmap = 'Greens',
            annot = True, fmt = "%.2f") # annot = True shows similarity values
plt.title('Product Similarity Heatmap - Top 10 Interacted Products')
plt.xticks(rotation = 45, ha = 'right')
plt.yticks(rotation = 0)
plt.show()
```



Let's look at the two closes ones with a similarity score of 0.35, which is far from perfectly similar. Product id #p067 and #p051.

products_df[product_interx['product_id'] == 'p067']

	product_id	name	category	strain_type	strain	vendor	THC_pct	CBD_pct	potency_mg	price_usd	availability	description	tags	date_added	text_feature
66	p067	Girl Scout Cookies	Flower	Sativa	Girl Scout Cookies	Sunset Wellness	24.0	1.3	NaN	61.49	410	Girl Scout Cookies Sativa flower with	['energetic', 'focused', 'uplifting']	2025-04-24	Girl Scout Cookies 3.5g Sativa Flower

products_df[product_interx['product_id'] == 'p051']

	product_id	name	category	strain_type	strain	vendor	THC_pct	CBD_pct	potency_mg	price_usd	availability	description	tags	date_added	text_feature
50	p051	OG Kush 1000mg	Concentrate	Sativa	OG Kush	Elevate Extracts	78.2	1.6	750.0	50.31	446	OG Kush Sativa concentrate with energetic, creative,	['energetic', 'creative', 'focused']	2025-03-14	OG Kush 1000mg Sativa Concentrate

7e. Top 5 Recommended Products Across All Users

```
from collections import Counter

# flatten all top-5 recommendations across all users
all_recommended = [pid for recs in recommendations.values() for pid in recs]

# count frequency of each product
top_counts = Counter(all_recommended)

# convert to dataframe for plotting
top_df = pd.DataFrame(top_counts.items(), columns = ['product_id','recommend_count'])
top_df = top_df.merge(products_df[['product_id','name']], on = 'product_id')
top_df = top_df.sort_values('recommend_count', ascending = False).head(5)

# plot
plt.figure(figsize = (8,5))
ax = sns.barplot(x = 'name', y = 'recommend_count', hue = 'name', data = top_df,
                palette = 'summer')

# remove grid lines
ax.grid(False)

# remove top and right spines
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

plt.title('Top 5 Most Recommended Products Across All Users')
plt.ylabel('Number of Recommendations')
plt.xlabel('Product Name')
plt.xticks(rotation = 45, ha = 'right')
plt.show()
```

